

Dynamic LoR-GS: Adaptive Low-Rank Gaussian Parameters for Memory-Efficient 3D Gaussian Splatting Training

[Viriyadhika Putra]

University of Toronto
Toronto, Canada
viriyadhika@cs.toronto.edu

[Aliza Tarakanov]

University of Toronto
Toronto, Canada
alizat@cs.toronto.edu

[Cosmo (Bo-Wei) Wen]

University of Toronto
Toronto, Canada
cosmo.wen@mail.utoronto.ca

Abstract

3D Gaussian Splatting (3DGS) achieves state-of-the-art novel view synthesis but suffers from prohibitive memory consumption, largely driven by the high-degree Spherical Harmonic (SH) coefficients required for view-dependent color. Current memory-reduction techniques, such as reducing SH degrees or applying uniform low-rank approximations (LoRA), apply uniform constraints to all primitives. This restriction degrades high-frequency specular details while wasting representational capacity on flat, diffuse geometry. In this paper, we present a Dynamic LoRA architecture that heterogeneously allocates parameter capacity across the scene based on physical complexity. By tracking representational demand via accumulated gradients, our system routes Gaussians into discrete rank bins using percentile-based quotas, scaling natively alongside geometric densification. This enables the network to automatically constrain capacity for matte surfaces while prioritizing complex, specular regions. Our evaluation shows that dynamic routing achieves perceptual parity (LPIPS) with full-capacity SH3 baselines at a fraction of the VRAM footprint. Furthermore, it consistently outperforms uniform Static LoRA models operating under the same effective parameter budget, demonstrating the advantage of scene-aware parameter distribution for 3DGS compression.

1 Introduction

The introduction of 3D Gaussian Splatting (3DGS) [8] has significantly advanced novel view synthesis by providing an explicit, point-based scene representation that achieves real-time rendering at high resolutions. However, this fidelity incurs a substantial computational cost. The explicit nature of 3DGS requires maintaining millions of independent primitives, resulting in a large GPU memory (VRAM) footprint during both training and inference. A majority of this parameter budget is consumed by the higher-order Spherical Harmonic (SH) coefficients, which are necessary to model complex, view-dependent color phenomena such as reflections and specular highlights.

To mitigate this bottleneck, recent works propose various compression strategies, ranging from geometric pruning [4, 19] to reducing the maximum allowed SH degree across the scene [2, 11, 14, 16]. Alternatively, applying standard Low-Rank Adaptation (LoRA) [6] to the SH parameters can compress the footprint. However, these static approaches share a fundamental limitation: they assume the scene’s view-dependent complexity is uniform. In physical reality, scenes are highly heterogeneous. A matte wall requires minimal view-dependent capacity, whereas a glossy car window requires precise, high-frequency parameterization. Imposing a uniform parameter constraint forces the over-parameterization of diffuse surfaces

while limiting the capacity available to render complex specularities accurately.

In this work, we address the limitations of uniform parameter allocation. We propose a Dynamic LoRA architecture that profiles and distributes representational capacity based on the localized demands of the scene. By factorizing the higher-order SH coefficients into a shared projection matrix and per-Gaussian low-rank latents, we create a flexible parameter space. During the geometric densification phase, we track the gradient activity of each primitive to quantify its demand for high-frequency details. Using a percentile-based routing mechanism to avoid spatial starvation, our system dynamically migrates Gaussians between discrete rank bins. This ensures that memory is allocated where the physical topology of the scene demands it.

In summary, our main contributions are:

- We introduce a Dynamic LoRA architecture for 3DGS that factorizes higher-order SH coefficients to enable heterogeneous, per-primitive capacity allocation.
- We develop a routing mechanism based on normalized gradient activity and percentile quotas, allocating higher representational capacity to complex primitives while preventing spatial starvation.
- We demonstrate empirically that our dynamic allocation achieves perceptual parity with full SH3 baselines at significantly reduced VRAM, and consistently outperforms uniform static baselines operating under identical effective parameter budgets.

2 Background and Related Work

3D Gaussian Splatting (3DGS) [8] is a leading framework for novel view synthesis. By representing scenes with explicit, differentiable 3D Gaussian primitives rendered via fast tile-based rasterization, 3DGS achieves high-fidelity real-time rendering and fast training. This efficiency has motivated extensive follow-on work in dynamic scenes [9], editing [3], compression [1, 11], and avatar generation [12].

Despite this success, optimizing 3DGS on memory-constrained hardware remains challenging due to two coupled factors. First, adaptive densification continuously clones and splits primitives, expanding the population N to several million. Second, each Gaussian requires a substantial parameter budget. Because view-dependent color is modeled by degree-3 spherical harmonic (SH3) coefficients, each primitive stores 48 floating-point values just for color. These coefficients dominate the memory footprint, consuming roughly 576 MB for a 3-million Gaussian scene before accounting for optimizer states and gradients [1, 11]. Consequently, peak training

VRAM is a critical bottleneck that existing methods struggle to fully resolve.

Current memory-reduction strategies generally fall into two categories: pruning and post-training compression. Pruning methods such as Mini-Splatting [4] and LP-3DGS [19] reduce N directly, which can degrade fine geometric detail. Alternatively, compression techniques like Reduced3DGS [11], LFGS [16], SG-Splatting [14], and MEGS2 [2] reduce the per-Gaussian footprint through SH band masking or lobe replacement. Because these operations occur *after* training completes, the full SH tensor must be maintained throughout optimization, leaving peak training memory unaffected.

We propose reducing the per-Gaussian parameter cost adaptively *during* training using Low-Rank Adaptation (LoRA) [6]. LoRA replaces a parameter matrix $W \in \mathbb{R}^{m \times n}$ with a low-rank factorization $W \approx AB$ (where $A \in \mathbb{R}^{m \times r}$, $B \in \mathbb{R}^{r \times n}$, and $r \ll \min(m, n)$). Maintaining only these low-rank factors directly reduces peak VRAM allocation.

The SH tensor $C \in \mathbb{R}^{3 \times 16}$ is the primary target for this factorization. Because methods like Reduced3DGS [11] successfully prune high-degree bands with minimal quality loss, we hypothesize the effective rank of C is inherently low for most primitives. Inspired by adaptive budget allocation in language models [17], we dynamically distribute this rank rather than applying a uniform constraint. This approach leverages the fact that matte surfaces require less representational capacity than complex specularities.

While LoRA has been introduced to the 3DGS ecosystem, it is typically applied to auxiliary neural network weights for dynamic streaming [9] or pose prediction [5]. StyleSplat [7] optimizes SH coefficients directly for style transfer but lacks a low-rank constraint. To our knowledge, this is the first work to apply low-rank factorization directly to 3DGS primitive parameters during optimization to reduce peak training memory.

3 Motivations and Objectives

While the peak training VRAM bottleneck of 3DGS [1, 8, 11] is well-documented, current mitigation strategies rely on uniform constraints that fail to account for the physical heterogeneity of real-world environments. Our primary motivation is grounded in the observation that view-dependent complexity varies significantly across a scene. A diffuse surface, such as a painted wall, exhibits minimal color variation across viewing angles and warrants only a few coefficients. Conversely, highly specular objects, such as a glossy car chassis, require precise, high-frequency parameterization spanning the full SH band. Assigning a uniform parameter budget to every Gaussian primitive is therefore structurally inefficient. Furthermore, while applying a static LoRA uniformly across the scene reduces total memory, it limits the capacity available for high-frequency details while over-parameterizing simple geometry, forcing an unnecessary compromise between VRAM savings and perceptual fidelity.

Driven by these observations, our core objective is to decouple a scene’s geometric resolution from its view-dependent memory footprint. We aim to design a dynamic, memory-aware routing architecture that autonomously distributes parameter capacity based on localized topological needs in the form of quantized low-rank buckets. By continuously evaluating the demand of each primitive

and actively readjusting its LoRA rank *during* the densification phase, this project seeks to demonstrate that adaptively allocating memory strictly to the regions that mathematically require it can significantly reduce peak training VRAM while preserving the high-frequency perceptual quality of the rendered scene.

4 Ideas and Methodologies

4.1 Isolating the Memory Bottleneck

For a scene with N Gaussian primitives, consider a target parameter block (e.g., the higher-order spherical harmonic coefficients) whose per-Gaussian contribution has dimension D . Instead of optimizing an unconstrained $[N \times D]$ parameter matrix, we decompose the *adjustment* to a frozen base using Low-Rank Adaptation (LoRA) [6] as:

$$\Delta = AB, \quad A \in \mathbb{R}^{N \times r}, \quad B \in \mathbb{R}^{r \times D}, \quad (1)$$

where $r \ll D$ is the rank. A is a per-Gaussian matrix, and B is a *shared* global projection matrix. The effective parameter used during rasterization is:

$$\theta_{\text{eff}} = \theta_{\text{base}} + \Delta, \quad (2)$$

where θ_{base} is frozen after initialization. A is optimized per-Gaussian, while B is optimized globally via a separate optimizer.

Alternatively, the full parameter matrix can be reconstructed directly as:

$$M = AB, \quad \theta_{\text{eff}} = AB, \quad (3)$$

provided the magnitude of θ_{eff} is sufficiently small.

Our initial implementation applies this parameterization with a fixed rank r to selected Gaussian attributes. The zero-order SH coefficient (mean color) remains a directly optimized parameter to preserve base color fidelity.

4.1.1 Flexible Target Selection. This architecture supports any subset of Gaussian attributes as LoRA targets. For example, jointly targeting colors (SH3, $D_c = 48$), scales ($D_s = 3$), and rotations ($D_q = 4$) yields a combined dimension $D = 55$. The shared matrix $B \in \mathbb{R}^{r \times 55}$ is partitioned column-wise as $B = [B_c \mid B_s \mid B_q]$. For Gaussian i with latent vector $a_i \in \mathbb{R}^r$, the effective attributes are:

$$\theta_{\text{color}}^{(i)} = \theta_{\text{color, base}}^{(i)} + a_i B_c, \quad (4)$$

$$\theta_{\text{scale}}^{(i)} = \theta_{\text{scale, base}}^{(i)} + a_i B_s, \quad (5)$$

$$\theta_{\text{quat}}^{(i)} = \theta_{\text{quat, base}}^{(i)} + a_i B_q. \quad (6)$$

A single latent vector a_i modulates all targeted attributes. We implement these updates as residuals to accommodate rotation and scale adjustments. Because frozen base tensors lack gradients, no Adam optimizer states are allocated for them.

4.1.2 LoRA-Aware Densification and Computation. Standard densification must be wrapped with LoRA-aware bookkeeping: cloning or splitting a Gaussian duplicates its corresponding row in A , while pruning removes it. During training, θ_{eff} is computed dynamically and passed to the rasterizer, ensuring regularization losses correctly penalize the effective parameters influencing the render.

4.1.3 Settling on Higher-Order SH. Ablation studies revealed that applying LoRA to geometry (positions, scales, rotations) degrades reconstruction quality, as these require precise, high-rank updates to represent sharp structural boundaries. Therefore, we restrict LoRA to higher-order SH bands (1–3). Because this tensor dominates the per-Gaussian parameter budget, it is simultaneously the most effective target for memory reduction.

4.2 Dynamic Resource Allocation

We iteratively designed and evaluated several dynamic allocation strategies to distribute capacity based on estimated primitive complexity.

4.2.1 Global Masking as a Proof of Concept. To validate the rasterizer’s ability to handle heterogeneous capacities, we first implemented a global masking strategy. We allocated a single high-rank tensor for the entire scene and simulated dynamic allocation by applying a binary mask during the forward pass to silence higher-order dimensions for selected Gaussians. While the approach is VRAM-inefficient due to the underlying dense allocation, this confirmed that heterogeneous rank representations maintain training stability, justifying the development of a genuinely sparse architecture.

4.2.2 Limitations of Absolute Scalar Thresholds. Our initial active migration strategy employed absolute scalar thresholds to promote or demote ranks based on gradient magnitudes. Empirical testing demonstrated this approach is brittle. Baseline gradient magnitudes fluctuate significantly across datasets depending on material properties (e.g., matte vs. reflective). Attempts to use statistical bounds, such as means or median multipliers, also failed. Looser bounds resulted in static allocations, while tighter bounds induced frequent rank thrashing as gradients oscillated during local error minimization. This establishes that absolute thresholding is ill-suited for the dynamic topology of 3DGS.

4.2.3 Gradient-Based Scoring and Percentile Routing. To quantify representational demand, we track gradient activity. For Gaussian i assigned to rank $r^{(i)}$, we accumulate the absolute gradients of its latent vector $a^{(i)}$ after each backward pass:

$$g^{(i)} += \|\nabla_{a^{(i)}} \mathcal{L}\|_1 \quad (7)$$

Post-epoch, this value is normalized by the current rank ($\tilde{g}^{(i)} = g^{(i)} / r^{(i)}$) to enable cross-bin comparison, and the buffer is reset.

To map these scores ($\tilde{g}$) to VRAM constraints, we compared fixed numerical quotas against percentile distributions. Theoretically, fixed quotas offer the minimal memory footprint: by establishing absolute bounds, the system can statically pre-allocate tensors, completely bypassing the VRAM overhead associated with dynamically slicing, copying, and migrating tensor entries. However, empirical implementation revealed a baseline “materialization cost”—a structural memory lower bound that both fixed and percentile routing inevitably encounter. Because this shared floor negates the theoretical memory advantage of static pre-allocation, we use percentile quotas ($q_{\text{top}}, q_{\text{mid}}, q_{\text{bot}}$). Although this requires dynamic tensor migration, percentiles scale natively with the evolving geometry, guaranteeing a consistent parameter-to-topology ratio.

4.2.4 Scene Analysis for Relative Thresholding. Because optimal quota fractions are scene-dependent, we require a mechanism to estimate them early in training, prior to reaching peak GPU memory. Inspired by post-training variance analysis [11], we evaluated two predictive signals during the densification phase: SH energy ($\|A_i B_i\|_F^2$) and initial color variance across training views:

$$s_i^{\text{color}} = \frac{1}{3} (\text{Var}(R_i) + \text{Var}(G_i) + \text{Var}(B_i)) \quad (8)$$

We evaluate the predictive accuracy of these signals against a step-15,000 full-SH3 gold standard, using GMM clustering to determine boundaries automatically.

4.2.5 Experimental Setup. All experiments were conducted on a single NVIDIA GeForce RTX 4070 Ti (12 GB VRAM) running Ubuntu 24.04.4 LTS with Python 3.12.13, PyTorch 2.11.0, CUDA 12.8, and gsplat 1.5.3. We evaluate on three scenes from the Ref-NeRF [13] real dataset (*gardenspheres*, *sedan*, and *toycar*), down-sampled by a factor of 4. Five model variants are compared: LoRA Complete, LoRA Fixed, LoRA Static, SH3, and SH2. All variants are trained for 15,000 steps with Gaussian densification stopped at step 10,000, a batch size of 1, and an SSIM loss weight of 0.2. To account for variance in Gaussian initialization and densification, each variant is run three times per scene, yielding 45 total training runs. Models are evaluated at steps 7,000 and 15,000 using PSNR, SSIM [15], and LPIPS [18] (AlexNet).

5 Results

To evaluate the computational and perceptual trade-offs of our architecture, we assess models across four primary metrics: Peak training VRAM (GB), PSNR, SSIM, and LPIPS (AlexNet). Because PSNR inherently favors over-smoothed pixel averages, we prioritize LPIPS as the primary indicator of preserved high-frequency, view-dependent specularities. Within our evaluation context, meaningful algorithmic improvements are defined as VRAM reductions of ≥ 0.5 GB, PSNR deltas of ≥ 0.1 dB, and LPIPS reductions of ≥ 0.01 .

5.1 Overall Performance and Memory Reduction

Table 1 presents the quantitative evaluation of our Dynamic LoRA architecture against the full-capacity Gold Standard (SH3) and a uniform Static LoRA baseline. Across all datasets, our dynamic routing achieves an optimal quality-to-memory tradeoff, significantly reducing the peak training VRAM footprint while preserving perceptual fidelity.

Compared to the SH3 baseline, our method achieves substantial memory savings: reducing peak VRAM by $\sim 26\%$ on the heavily populated *Garden Sphere* scene ($7.08 \rightarrow 5.18$ GB) and up to 36% on the *Sedan* scene ($2.90 \rightarrow 1.85$ GB). Notably, despite the bookkeeping overhead of maintaining dynamic routing indices, our method achieves a memory footprint that is either highly competitive with or strictly superior to the uniform Static LoRA baseline (e.g., achieving the lowest footprint on the *Sedan* dataset).

Importantly, this VRAM reduction does not incur a proportional penalty to rendering quality. As shown in Table 1, Dynamic LoRA consistently secures the “best” or “second-best” position across most quantitative metrics. In terms of PSNR and SSIM, our architecture

Dataset	Method	Peak VRAM (GB) ↓	PSNR (db) ↑	SSIM ↑	LPIPS ↓
Garden Sphere	SH3 (Gold Standard)	7.08 ± 0.00	22.899 ± 0.046	0.614 ± 0.000	0.219 ± 0.001
	Static LoRA	4.74 ± 0.01	22.865 ± 0.001	0.614 ± 0.000	0.225 ± 0.000
	Dynamic LoRA (Ours)	5.18 ± 0.01	22.875 ± 0.009	0.612 ± 0.000	0.217 ± 0.000
Sedan	SH3 (Gold Standard)	2.90 ± 0.01	25.430 ± 0.019	0.719 ± 0.000	0.274 ± 0.000
	Static LoRA	1.96 ± 0.00	25.443 ± 0.030	0.720 ± 0.000	0.274 ± 0.000
	Dynamic LoRA (Ours)	1.85 ± 0.02	25.427 ± 0.016	0.718 ± 0.000	0.276 ± 0.000
Toycar	SH3 (Gold Standard)	3.42 ± 0.00	24.819 ± 0.005	0.678 ± 0.000	0.192 ± 0.001
	Static LoRA	2.40 ± 0.01	24.663 ± 0.002	0.673 ± 0.000	0.196 ± 0.000
	Dynamic LoRA (Ours)	2.41 ± 0.01	24.771 ± 0.007	0.677 ± 0.000	0.194 ± 0.001

Table 1: Overall Performance and Memory Reduction. Bold indicates the best result, and underline indicates the second-best per dataset.



Figure 1: Qualitative Evaluation and Dynamic Rank Distribution on the *Garden Sphere* Scene. (Left to right: Ground Truth, SH3 Baseline, Dynamic LoRA (Ours), Dynamic Rank Heatmap). Our Dynamic LoRA architecture achieves perceptual parity with the unconstrained SH3 baseline, preserving high-frequency details and specular reflections. The rank heatmap visualizes our dynamic allocation strategy: the system constrains diffuse backgrounds to a minimal capacity ($r_{\min}$, blue), allocates moderate capacity to structural edges (r_{mid} , purple), and concentrates its maximum parameter budget on complex, view-dependent specularities ($r_{\max}$, red). This demonstrates that our routing mechanism actively learns and exploits the physical topology of the scene.

closely matches the unconstrained SH3 baseline, demonstrating that dynamically constraining parameters does not destabilize geometric convergence. Furthermore, on the highly complex *Garden Sphere* dataset, our dynamic method outperforms the SH3 Gold Standard in LPIPS (0.217 vs. 0.219). This validates our core architectural hypothesis: by adaptively allocating parameter capacity strictly to regions that demand it, the network is forced to learn robust structural representations rather than overfitting high-degree Spherical Harmonics to noise, thereby maintaining sharp, view-dependent realism at a fraction of the hardware cost.

5.2 Attribute Ablation for LoRA Targeting

We evaluate the effect of LoRA target selection on reconstruction quality (PSNR) across three scenes (*Garden Sphere*, *Sedan*, *Toycar*). The full-rank SH3 baseline operates without any LoRA constraint.

As shown in Table 2, all LoRA configurations remain competitive with the full-rank baseline, maintaining quality gaps generally under 0.5 dB. Isolating LoRA strictly to colors achieves the closest match to the baseline on the *Sedan* scene (25.14 vs. 25.16 dB) and performs consistently across all scenes. Expanding the target to include rotations yields nearly identical results, suggesting rotations

possess an inherent low-rank structure that LoRA captures without significant loss.

Conversely, including scales in the target set introduces the most notable degradation. Configurations targeting scales drop by approximately 0.75 dB on *Sedan* compared to the colors-only approach, indicating that scale parameters require full-rank updates to preserve fine geometric detail. Furthermore, targeting solely geometry attributes (rotations and scales) achieves competitive PSNR on *Garden Sphere* and *Toycar* but underperforms on *Sedan*, highlighting a scene-dependent sensitivity to geometric factorization.

Overall, targeting only colors provides the most predictable and robust behavior across varying topologies, justifying its selection as the exclusive target for our dynamic rank architecture.

5.3 Evaluating Heuristics for Scene-Aware Partitioning

Across all combinations of scoring signal and clustering method, rendering quality differences were negligible (PSNR range 24.33–24.44 dB, standard deviation ~ 1.0 dB across three scenes). This suggests that the choice of per-Gaussian scoring signal is less critical than the overall quota distribution.




Figure 2: Gold-standard quota fractions (SH3 at step 15,000, GMM, color variance) across seven scenes.

Target	Garden Sphere	Sedan	Toy car
Full-rank SH3 (Baseline)	23.29	25.16	25.01
Colors Only (SH 1–3)	23.08	25.14	24.85
Colors + Rotations	23.10	25.02	24.87
Colors + Opacities	22.95	25.05	24.78
Colors + Scales	22.94	24.40	24.81
Colors + Rot. + Scales	22.94	24.44	24.81
Rot. + Scales (Ref.)	23.14	24.30	24.98
Rot. + Scales + Opac. (Ref.)	23.12	24.35	25.05

Table 2: Ablation of LoRA target configurations on reconstruction quality (PSNR (db)). Bold indicates the best result among color-targeting configurations per scene. Geometry-only rows serve as baselines and are not candidates for dynamic memory reduction.

Figure 2 shows gold-standard quota fractions across seven scenes using color variance. These fractions vary substantially with the low-rank bucket ranging from 0.42 to 0.59, confirming that scene-independent fixed quotas are ineffective and justifying scene-specific estimation.

Figure 3 compares prediction error over training steps for color variance and SH energy, evaluated under rank-2 and rank-32 warmups. Because color variance depends solely on Gaussian geometry, it is unaffected by warmup rank and consistently outperforms SH energy. Conversely, SH energy begins with large errors because the LoRA factors lack meaningful initialization. While the rank-32 model recovers by step 2100, the rank-2 model remains volatile until step 6000, and both remain less accurate than color variance. We therefore train at rank 2 until step 4000 and run a single color




Figure 3: Mean absolute error vs. gold standard (GMM, averaged across scenes). Color variance (blue/green) remains below SH energy (orange/red) at all steps, including step 0.

variance pass to establish scene-specific quotas. This strategy minimizes memory consumption through the densification phase while providing a reliable quota estimate before reaching peak memory.

5.4 Effective Rank Analysis and Architectural Trade-offs

To isolate the specific impact of our dynamic routing mechanism, we evaluate our architecture against a uniformly allocated Static LoRA baseline and a standard reduced-capacity SH2 baseline. To ensure a fair comparison, all models operate under a roughly equivalent parameter budget: SH2 natively utilizes 15 parameters per channel, Static LoRA is fixed uniformly at $r = 16$, and our Dynamic LoRA averages an effective rank of $r \approx 16.4$ across the population.

By architectural definition, standard SH2 establishes the lower bound for memory consumption (e.g., 4.50 GB on the *Garden Sphere* scene). Because LoRA-based methods require maintaining localized pointer arrays, bucket indices, and segregated Adam optimizer states, they inherently carry a VRAM overhead that prevents them from outperforming SH2 in pure memory reduction at the same effective rank.

However, evaluating these representations strictly through global quantitative metrics (PSNR and SSIM) obscures the critical advantage of heterogeneous allocation. While all methods achieve effective parity in these global averages, spatial averaging dilutes localized perceptual improvements. Standard SH2 and Static LoRA constrain every primitive to a uniform capacity, limiting the reconstruction of complex, high-frequency specularities. Consequently, while they save memory, they tend to over-smooth fine details.

In contrast, our dynamic architecture leverages the same average parameter budget but distributes it heterogeneously. As demonstrated by the LPIPS metric, which heavily penalizes the loss of

Dataset	Method	Peak VRAM (GB) ↓	PSNR (db) ↑	SSIM ↑	LPIPS ↓
<i>Garden Sphere</i>	SH2 (Low-Mem Baseline)	4.50 ± 0.00	22.938 ± 0.009	0.616 ± 0.000	0.225 ± 0.001
	Fixed (Dynamic Rank Only)	5.77 ± 0.01	<u>22.881 ± 0.005</u>	0.612 ± 0.000	0.217 ± 0.001
	Static (Uniform Rank)	<u>4.74 ± 0.01</u>	22.865 ± 0.001	<u>0.614 ± 0.000</u>	<u>0.225 ± 0.000</u>
<i>Sedan</i>	SH2 (Low-Mem Baseline)	1.85 ± 0.02	25.328 ± 0.018	0.716 ± 0.000	0.277 ± 0.000
	Fixed (Dynamic Rank Only)	2.35 ± 0.01	<u>25.431 ± 0.013</u>	<u>0.719 ± 0.000</u>	<u>0.274 ± 0.001</u>
	Static (Uniform Rank)	<u>1.96 ± 0.00</u>	25.443 ± 0.030	0.720 ± 0.000	0.274 ± 0.000
<i>Toyicar</i>	SH2 (Low-Mem Baseline)	2.23 ± 0.00	24.805 ± 0.008	0.680 ± 0.000	0.196 ± 0.000
	Fixed (Dynamic Rank Only)	2.86 ± 0.01	<u>24.750 ± 0.006</u>	0.675 ± 0.000	0.193 ± 0.000
	Static (Uniform Rank)	<u>2.40 ± 0.01</u>	24.663 ± 0.002	0.673 ± 0.000	<u>0.196 ± 0.000</u>

Table 3: Ablation of effective rank and architectural trade-offs. Bold indicates the best result, and underline indicates the second-best per dataset.

sharp, view-dependent textures, dynamic routing consistently outperforms both SH2 and Static LoRA on complex scenes. By constraining flat, diffuse geometry that requires minimal capacity, the system frees parameters to prioritize complex specular regions. This demonstrates that dynamic rank distribution is structurally superior to a static uniform rank: it matches the physical topology of the scene, achieving higher perceptual realism without requiring a larger global parameter budget.

6 Discussion

6.1 Active Allocation vs. Post-Hoc Pruning

Mallick et al. [10] address peak training VRAM by replacing unconstrained densification with a budget-constrained constructive process. While this systematically degrades quality as Gaussian counts (N) shrink, it primarily targets geometric redundancy. Our work is fundamentally orthogonal, targeting representation redundancy within the SH coefficients. Because high spatial frequency does not strictly correlate with high view-dependency, these approaches are complementary. As Papantonakis et al. [11] establish that primitive count and SH storage are the two dominant memory bottlenecks in 3DGS, a hybrid approach is highly promising. Combining budget-constrained densification with our dynamic LoRA rank assignment would enable a dual-budget optimization framework. Such a system could allocate resources precisely by spending the “Gaussian budget” on geometric detail and the “Rank budget” on complex view-dependent effects. This approach would likely achieve a more favorable Pareto frontier than either method in isolation.

6.2 Material-Driven Limits and the Oracle Problem

A critical limitation arises in globally uniform environments (e.g., the matte *Toyicar* dataset), where perfectly aligned native representations (e.g., universal SH1) outperform our dynamic architecture by avoiding tensor slicing and routing overhead. This exposes an oracle problem in memory-efficient 3DGS: mathematically optimizing the percentile quotas (q_{top} , q_{bot}) requires prior knowledge of the scene’s final material complexity. However, because geometry is volatile during early densification, this complexity is only revealed

by training. This creates a functional impasse: perfectly configuring the memory-saving architecture requires a full training run, which inherently renders the architecture itself redundant. Resolving this paradox necessitates future work into pre-optimization estimators capable of predicting scene complexity directly from the initial sparse point cloud.

6.3 Implementation Overhead and Future Integration

Despite requiring per-bucket optimizers and periodic reallocation, our implementation achieves memory savings close to the theoretical limit. For the *Garden Sphere* scene, our average rank of 16.4 predicts a theoretical saving of 1.48 GB against a measured 1.31 GB. This 0.17 GB gap (similarly minimal in *Sedan* and *Toyicar*) is achieved through strict Python-level memory management, explicitly clearing stale optimizer states during rank reallocation. The remaining discrepancy represents a framework-level materialization cost. Specifically, temporary dense tensors must be allocated during densification and forward passes to interface fragmented bucket matrices with standard framework logic. To completely eliminate this indexing overhead and maximize hardware efficiency, variable-rank memory pointers must eventually be integrated into custom C++/CUDA kernels. Nevertheless, our implementation validates the architectural viability of dynamic capacity routing.

7 Conclusion

We demonstrated that the prohibitive memory footprint of 3DGS is primarily a consequence of uniform parameter allocation. Our Dynamic LoRA architecture employs gradient-based percentile routing to identify and exploit heterogeneous representational demands, automatically constraining diffuse geometry to prioritize capacity for complex specularities. Consequently, our method achieves the perceptual realism of a full-capacity SH3 model at a fraction of the memory cost. While current framework-level tensor management imposes a minor materialization overhead, the architectural viability of dynamic routing motivates the next critical step: integrating variable-rank parameter resolution directly into custom CUDA kernels. By aligning mathematical capacity with physical scene topology, we provide a scalable foundation for memory-efficient 3DGS architectures.

References

- [1] Milena T. Bagdasarian, Paul Knoll, Yi-Hsin Li, Florian Barthel, Anna Hilsmann, Peter Eisert, and Wieland Morgenstern. 2025. 3DGS.zip: A survey on 3D Gaussian Splatting Compression Methods. *Comput. Graph. Forum* 44, 2 (2025). doi:10.1111/CGF.70078
- [2] Jiarui Chen, Yikeng Chen, Yingshuang Zou, Ye Huang, Peng Wang, Yuan Liu, Yujing Sun, and Wenping Wang. 2025. MEGS²: Memory-Efficient Gaussian Splatting via Spherical Gaussians and Unified Pruning. *CoRR* abs/2509.07021 (2025). arXiv:2509.07021 doi:10.48550/ARXIV.2509.07021
- [3] Yiwen Chen, Zilong Chen, Chi Zhang, Feng Wang, Xiaofeng Yang, Yikai Wang, Zhongang Cai, Lei Yang, Huaping Liu, and Guosheng Lin. 2024. GaussianEditor: Swift and Controllable 3D Editing with Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2024, Seattle, WA, USA, June 16-22, 2024*. IEEE, 21476–21485. doi:10.1109/CVPR52733.2024.02029
- [4] Guangchi Fang and Bing Wang. 2024. Mini-Splatting: Representing Scenes with a Constrained Number of Gaussians. In *Computer Vision - ECCV 2024 - 18th European Conference, Milan, Italy, September 29-October 4, 2024, Proceedings, Part LXXXVII (Lecture Notes in Computer Science)*, Ales Leonardis, Elisa Ricci, Stefan Roth, Olga Russakovsky, Torsten Sattler, and Gül Varol (Eds.). Springer, 165–181. doi:10.1007/978-3-031-72980-5_10
- [5] Yuki Fujimura, Takahiro Kushida, Kazuya Kitano, Takuya Funatomi, and Yasuhiro Mukaigawa. 2025. UVF-Splatter: Pose-Free Feed-Forward 3D Gaussian Splatting Adapted to Unfavorable Views. *CoRR* abs/2507.22342 (2025). arXiv:2507.22342 doi:10.48550/ARXIV.2507.22342
- [6] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net. <https://openreview.net/forum?id=nZeVKeeFYf9>
- [7] Sahil Jain, Avik Kuthiala, Prabhdeep Singh Sethi, and Prakanshul Saxena. 2024. StyleSplat: 3D Object Style Transfer with Gaussian Splatting. *CoRR* abs/2407.09473 (2024). arXiv:2407.09473 doi:10.48550/ARXIV.2407.09473
- [8] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 2023. 3D Gaussian Splatting for Real-Time Radiance Field Rendering. *ACM Trans. Graph.* 42, 4 (2023), 139:1–139:14. doi:10.1145/3592433
- [9] Zhenhuan Liu, Shuai Liu, Yidong Lu, Yirui Chen, Jie Yang, and Wei Liu. 2025. Efficient 4D Gaussian Stream with Low Rank Adaptation. *CoRR* abs/2502.16575 (2025). arXiv:2502.16575 doi:10.48550/ARXIV.2502.16575
- [10] Saswat Subhrajyoti Mallick, Rahul Goel, Bernhard Kerbl, Markus Steinberger, Francisco Vicente Carrasco, and Fernando De la Torre. 2024. Taming 3DGS: High-Quality Radiance Fields with Limited Resources. In *SIGGRAPH Asia 2024 Conference Papers, SA 2024, Tokyo, Japan, December 3-6, 2024*, Takeo Igarashi, Ariel Shamir, and Hao (Richard) Zhang (Eds.). ACM, 2:1–2:11. doi:10.1145/3680528.3687694
- [11] Panagiotis Papantonakis, Georgios Kopanas, Bernhard Kerbl, Alexandre Lanvin, and George Drettakis. 2024. Reducing the Memory Footprint of 3D Gaussian Splatting. *Proc. ACM Comput. Graph. Interact. Tech.* 7, 1 (2024), 16:1–16:17. doi:10.1145/3651282
- [12] Shenhan Qian, Tobias Kirschstein, Liam Schoneveld, Davide Davoli, Simon Giebenhain, and Matthias Nießner. 2024. GaussianAvatars: Photorealistic Head Avatars with Rigged 3D Gaussians. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2024, Seattle, WA, USA, June 16-22, 2024*. IEEE, 20299–20309. doi:10.1109/CVPR52733.2024.01919
- [13] Dor Verbin, Peter Hedman, Ben Mildenhall, Todd E. Zickler, Jonathan T. Barron, and Pratul P. Srinivasan. 2022. Ref-NeRF: Structured View-Dependent Appearance for Neural Radiance Fields. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*. IEEE, 5481–5490. doi:10.1109/CVPR52688.2022.00541
- [14] Yiwen Wang, Siyuan Chen, and Ran Yi. 2025. SG-Splatting: Accelerating 3D Gaussian Splatting with Spherical Gaussians. *CoRR* abs/2501.00342 (2025). arXiv:2501.00342 doi:10.48550/ARXIV.2501.00342
- [15] Zhou Wang, Alan C. Bovik, Hamid R. Sheikh, and Eero P. Simoncelli. 2004. Image quality assessment: from error visibility to structural similarity. *IEEE Trans. Image Process.* 13, 4 (2004), 600–612. doi:10.1109/TIP.2003.819861
- [16] Ruiqi Wu, Linlin Jiao, Gang Liu, Li Zhu, Xuan Fei, Yashuang Mu, and Chao Fan. 2025. LFGS: A lightweight framework for efficient 3D Gaussian Splatting with minimal memory footprint. *Comput. Graph.* 131 (2025), 104309. doi:10.1016/J.CAG.2025.104309
- [17] Qingru Zhang, Minshuo Chen, Alexander Bukharin, Pengcheng He, Yu Cheng, Weizhu Chen, and Tuo Zhao. 2023. Adaptive Budget Allocation for Parameter-Efficient Fine-Tuning. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net. <https://openreview.net/forum?id=lq62uWRJjiY>
- [18] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. 2018. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. In *2018 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2018, Salt Lake City, UT, USA, June 18-22, 2018*. Computer Vision Foundation / IEEE Computer Society, 586–595. doi:10.1109/CVPR.2018.00068
- [19] Zhaoliang Zhang, Tianchen Song, Yongjae Lee, Li Yang, Cheng Peng, Rama Chellappa, and Deliang Fan. 2024. LP-3DGS: Learning to Prune 3D Gaussian Splatting. In *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*, Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang (Eds.). http://papers.nips.cc/paper_files/paper/2024/hash/dd51dbce305433cd60910dc5b0147be4-Abstract-Conference.html